

UNIVERSITÀ  
DEGLI STUDI  
DI PADOVA

# Giacomo Gasparotto **Assignment 1**

Quantum Information and Computing  
A.Y. 25/26  
11 November 2025



# Integer overflow

Sum **2000000** and **1** in two settings:

**int16**:  $-31616 + 1 = -31615$

**Why?**

- Let's look at the two-complement:

$$(2000000)_{10} = (111\ 110\ 100\ 001\ 001\ 000\ 000)_2$$

- If we store the number in int16 it takes the first 16 most significant digits:

$$(1110\ 0010\ 0100\ 0000)_2 = (-31616)_{10}$$

Leading to an error of overflow!

**int32**:  $2000000 + 1 = 2000001$

Enough space to store the result.

Sum of 2000000 and 1 using 16-bit integers:  
 $-31616 + 1 = -31615$

Max interval for 16-bit: [-32768, 32767]

Sum of 2000000 and 1 using 32-bit integers:  
 $2000000 + 1 = 2000001$

Max interval for 32-bit: [-2147483648, 2147483647]



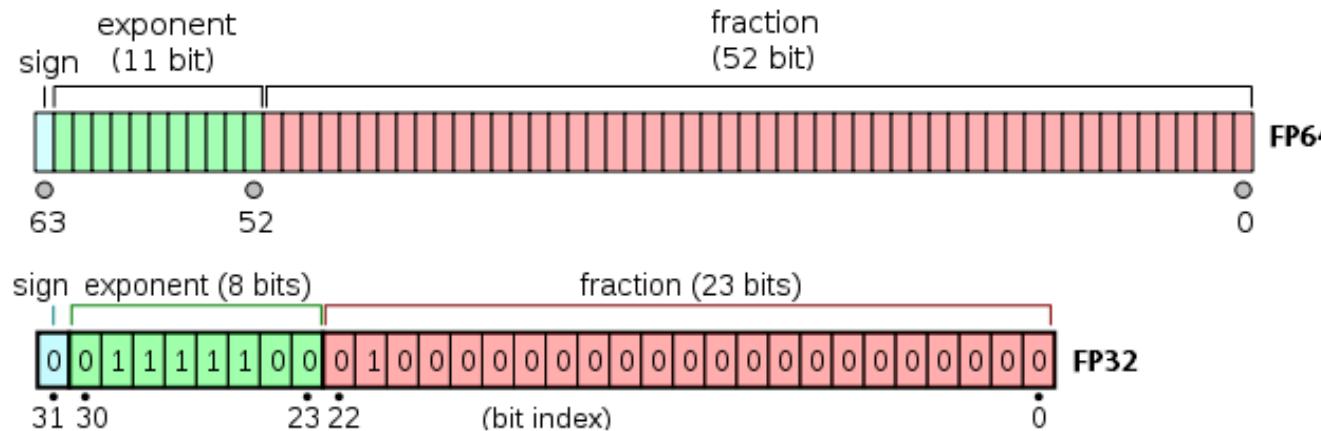
# Floating point rounding error

Sum two large real numbers in two settings:

**real32** – single precision

- $\pi \cdot 10^{32}$ : 3.14159278E+32
- $\sqrt{2} \cdot 10^{21}$ : 1.41421360E+21
- **sum**: 3.14159278E+32

The second number is “lost” due to **rounding** because of 10 order of magnitudes of difference.



**real64** – double precision

- $\pi \cdot 10^{32}$ : 3.1415926535897933E+032
- $\sqrt{2} \cdot 10^{21}$ : 1.4142135623730950E+021
- **sum**: 3.1415926536039354E+03

The mantissa is large enough to store the entire sum.

```
Single precision real numbers to sum:  
3.14159278E+32  
1.41421360E+21  
Single precision sum: 3.14159278E+32  
  
Double precision real numbers to sum:  
3.1415926535897933E+032  
1.4142135623730950E+021  
Double precision sum: 3.1415926536039354E+032
```



# Matrix multiplication

Method	CPU usage	Performance
<b>row-by-col</b>	Contiguous	Efficient cache usage
<b>col-by-row</b>	Non contiguous	Poor cache locality
<b>matmul</b>	Contiguous	Optimized (BLAS subroutines)

- Fortran uses **column-major order** to store matrices, that's why the second method is expected to be less efficient in terms of CPU than the other two

## CPU performance:

- CPU time check
- Optimization flags (-O2, -O3): optimize matrix loops by reordering, unrolling and parallelizing them

Column-major order

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$



# Code development

Matrix-matrix multiplication loops:

```
! Matrix multiplication (row-by-column)
do i = 1, m
    do j = 1, p
        C(i,j) = 0.0_real32
        do k = 1, n
            C(i,j) = C(i,j) + A(i,k) * B(k,j)
        end do
    end do
end do
```

```
! Matrix multiplication (column-by-row)
do j = 1, p
    do i = 1, m
        C(i,j) = 0.0_real32
        do k = 1, n
            C(i,j) = C(i,j) + A(i,k) * B(k,j)
        end do
    end do
end do
```

CPU-time loop:

```
open(unit=10, file='row_col.dat', status="unknown", position="append")
call cpu_time(init_time)
call matmul_rowbycol(M1, M2, prod1, dim, dim, dim, .false.)
call cpu_time(end_time)
write(10,*) dim, end_time - init_time
close(10)

open(unit=10, file='col_row.dat', status="unknown", position="append")
call cpu_time(init_time)
call matmul_colbyrow(M1, M2, prod2, dim, dim, dim, .false.)
call cpu_time(end_time)
write(10,*) dim, end_time - init_time
close(10)

open(unit=10, file='matmul.dat', status="unknown", position="append")
call cpu_time(init_time)
prod3 = matmul(M1, M2)
call cpu_time(end_time)
write(10,*) dim, end_time - init_time
close(10)
```



# Results

